

The LPARDesign and LPARDesign Extended Tools

Alain Maneville
Executive I/T Specialist
zCA, zChampion

Thierry Deleris
Responsable de Mission Transverse
BPCE-IT

17 Juin 2021



IBM Z

AGENDA

✓ PART-1

- ☐ Background of HiperDispatch® - Purpose of the Tool – Why the Tool ?
- ☐ HiperDispatch® Overview
- ☐ LPARDesign Tool demo and explanations

✓ PART-2

- ☐ LPARDesign Extended Tool

PURPOSE OF THE TOOL – WHY THE TOOL

✓ Background of HiperDispatch®

- HiperDispatch® was announced with the z10
- No actual IBM presentation / actual explanation / Tool for this feature were available
 - The z10 HiperDispatch® presentation was funny

✓ Purpose of the Tool

- Ease the LPAR configuration
- Provide optimizations of the LPAR configuration
- Provide visualization of the LPAR configuration

✓ Why the Tool?

- HiperDispatch® needs some calculations to have the spread of the Vertical LPs
- The calculations are sometimes different depending on the machine type
- The Tool makes all the calculations for you
- The Tool can “talk” to zPCR, so can be part of the IBM sizing methodology.
 - *You can generate a zPCR Study*
 - *You can generate an .XML configuration definition to be read by LPARDesign Extended*

HiperDispatch® Overview

Hiperdispatch®: Motivation

- **Design Objective**
 - **Keep work** as much as possible **local to a physical processor** to optimize the usage of the processor caches
 - **Expected Result**
 - Cache reloads shall occur much less often
 - Cache misses and fetches from other books (and chips) should be avoided as much as possible
- **Function: Hiperdispatch®**
 - Interaction between z/OS and PR/SM to optimize work unit and logical processor placement to physical processors
 - Consists of 2 parts
 - In z/OS (sometimes referred as **Dispatcher Affinity**)
 - Because it attempts to create a temporary affinity between work and processors
 - In PR/SM (sometimes referred as **Vertical CPU Management**)
 - Because it attempts to assign physical processors exclusively to logical processors (as much as possible)



■ OPTIMIZING THE NUMBER OF LOGICAL PROCESSORS DOWN TO THE MINIMUM NUMBER NEEDED OF PHYSICAL PROCESSORS

–Based on the Share of the logical partition

1 – Compute the %Share of the LPAR

$$\%Share(LPARi) = \frac{Weight(LPARi)}{\sum_{j=1}^n Weight(LPARj)}$$

2 – Compute the Guaranteed number of PP this LPAR is entitled to

$$\#PP(LPARi) = \%Share(LPARi) * Total_ \#PP_of_Shared_Pool = \text{N.M}$$

• Compute Result = N.M → Number of PP guaranteed to the LPAR

- **N** = number of physical processors which can be **used completely** by this partition
- **M** = the **fraction of a physical processor** which must be used to satisfy the share of the partition

• Example – A Machine with 5 PP

- Assignment of logical processors to physical processors in HiperDispatch® mode
- LPAR1 – **N=3**, **M=0.5**
 - 3 physical processors (High Processors)
 - Share of 50% of the 4th processor (Medium Processor)
- LPAR2 – **N=1**, **M=0.5**
 - 1 physical processor (High Processors)
 - Share of 50% of the 2nd processor

Partition	LPs	Weight	Share	Share in PPs
LPAR1	5	350	70%	3.5
LPAR2	5	150	30%	1.5
		500		5

• What about the “others” Logical Processors?

- **1** for LPAR1 and **3** for LPAR2
 - Low Processors (parked = not used)
- If demand exists AND the other partition does not need their share
 - Medium processors can use up to all of their physical processors
 - Low processors can be un-parked and start to use physical processors which are not needed by other partitions

New wording with HiperDispatch®

Vertical High **VH**

- uses 100% of a PP

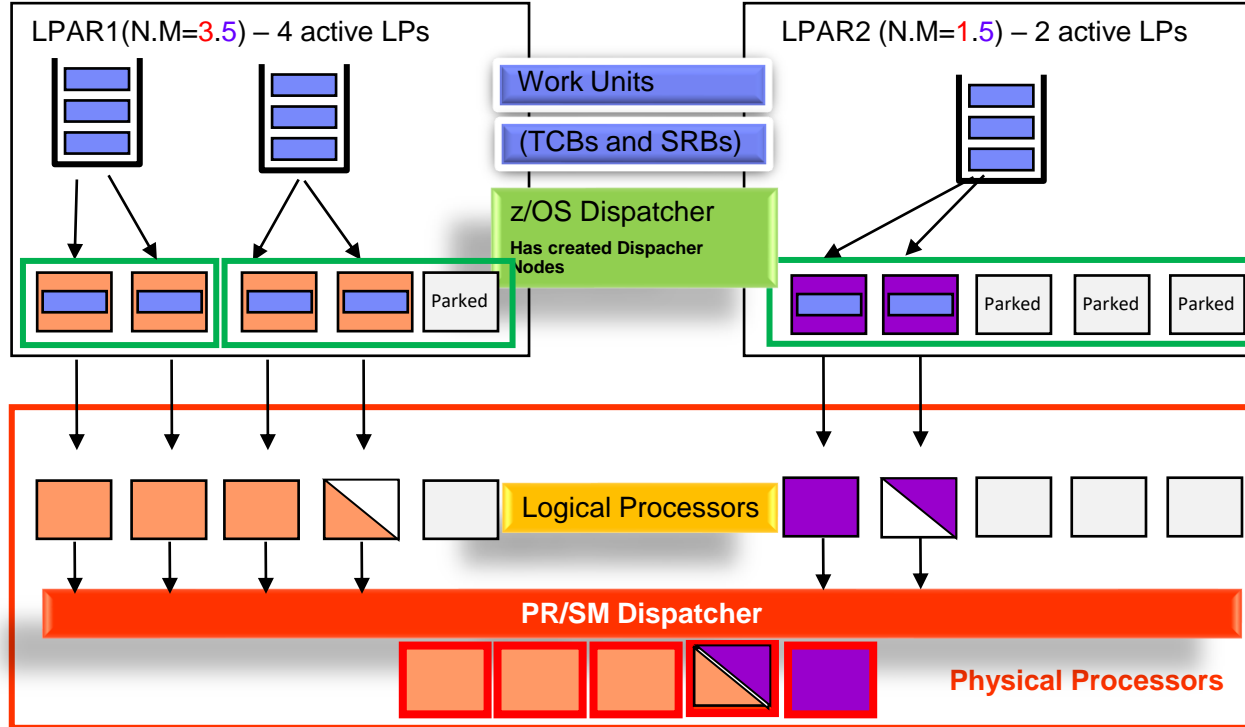
Vertical Medium **VM**

- uses 1-99% of a PP (100% in specific cases)

Vertical Low **VL**










- can be “Unparked” if demand exists and thus joins the VM pool

Dispatching in a z/OS and PR/SM environment (w/HD)



Hiperdispatch : PR/SM® - What has changed in z14 and z15 vs z13

- LP Spread Computing

Case#	Specify Machine from the list	Type	Specify N.M	Specify #LP		#VH	#VM	VM%	#VL	#Active	Comments
1	8561	z15	1.50	2		 1	1	50.00%	 0	2	
2	3906	z14	1.50	2		 1	1	50.00%	 0	2	
3	2964	z13	1.50	2		 0	2	75.00%	 0	2	z13,z13s Rule even if M>=0.5

The HiperDispatch® formula

The Basic HiperDispatch® Equation

Compute Result = $N.M$ → Number of Physical Procs. guaranteed to the LPAR

- **N** = number of physical processors which can be **used completely** by this LPAR
- **M** = the **fraction of a physical processor** which must be used to satisfy the share of the LPAR

$$N.M = \#PP * \%SHARE$$

Number of Physical Processor
May change after:

- ON/OFF
- Adding an LPAR with DED Procs.

%Share – May change after:

- Adding a new LPAR
- Deleting an LPAR
- Changing the Weight of a LPAR
- Soft Capping with positive Phantom Weight

Hiperdispatch® : PR/SM part .. – LP entitlements – #LP rules

Do not worry !!!!! All these rules and calculations are done for you in the **LPARDesign Tool**


ID=IBM Corp. - LPARDesign V11-T03 Current zPCR Version-9.4b - SpecCfg=YES

This macro evaluates the LPAR definition for HD eligible processors

To create a copy of this spreadsheet

Create a copy...

Optional but recommended



To start

Accept SPECIAL Conf ?
YES

Change Customer Name / ID

2 - Change Configuration

Machine-Type	8561-718	==> Not HD Eligible
#zIIP	18	
#ICF	18	
#IFL	18	

To save results =====>

Save as...

© Copyright IBM France 2020

DISCLAIMER OF WARRANTIES:

The following [enclosed] macro is sample code created by **Alain Maneville - IBM France**.
 This sample macro is not part of any standard IBM product and is provided to you solely for the purpose of assisting you in the PR/SM LPAR Configuration.
 The code is provided "AS IS", without warranty of any kind. IBM shall not be liable for any damages arising out of your use of such sample code, even if you have been advised of the possibility of such damage.

The DASHBOARD-HD code and the zPCR code have been written by **Thierry Deleris (a customer from BPCE-IT)**
 The CONFIG-MSU code has been written by **Robert Vaupel - IBM (STSM, z/OS Workload and Performance Management)**
 Support : Support will be provided on a "best effort" basis. Send the spreadsheet for an analysis to alain_maneville@fr.ibm.com

LparDesign
✕

Minimize/Maximize

Go To Start

GCP

zIIP

IFL

ICF

✓

i

🔍

📊

✓

i

🔍

📊

✓

i

🔍

📊

✓

i

🔍

📊

Go To MSU

Go To Dashboard

Go To Expert

zPCR Menu

XML Menu

Print

Configuration

Model
8561-718
☐

zIIP
18
☐

IFL
18
☐

ICF
18
☐

First Worksheet – Defining the configuration, and some parameters.

➡ ID=IBM Corp. - LPARDesign V11-T03 Current zPCR Version-9.4b - SpecCfg=YES

This macro evaluates the LPAR definition for HD eligible processors

To create a copy of this spreadsheet

Create a copy...

Optional but recommended

PR/SM accepts configurations where the number of the HMC defined LPs is not consistent with the number of LPs needed to sustain the share of the LPAR

To start

Accept SPECIAL Conf ? **YES**

Change Customer Name / ID

} The Study ID

2 - Change Configuration	
Machine-Type	8561-718
#zIIP	18
#ICF	18
#IFL	18

} The configuration

==> Not HD Eligible

To save results

=====>

Save as...

© Copyright IBM France 2020

DISCLAIMER OF WARRANTIES:

The following [enclosed] macro is sample code created by **Alain Maneville - IBM France**.

This sample macro is not part of any standard IBM product and is provided to you solely for the purpose of assisting you in the PR/SM LPAR Configuration. The code is provided "AS IS", without warranty of any kind. IBM shall not be liable for any damages arising out of your use of such sample code, even if you have been advised of the possibility of such damage.

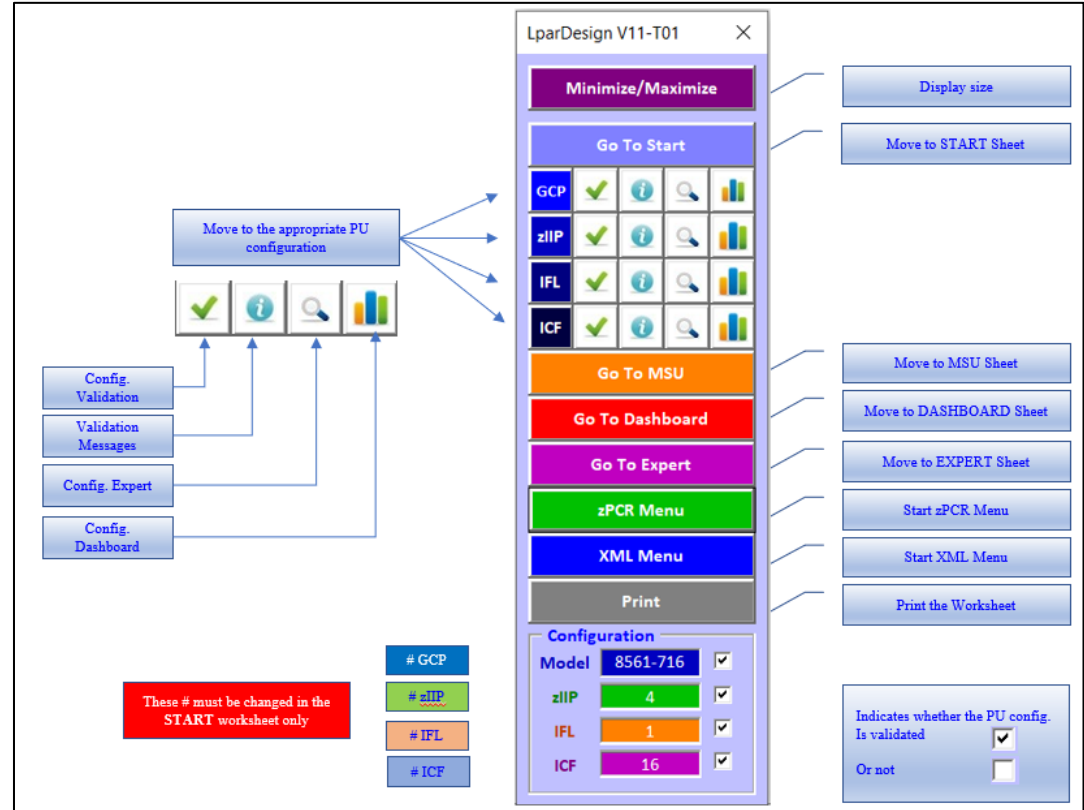
The DASHBOARD-HD code and the zPCR code have been written by **Thierry Deleris (a customer from BPCE-IT)**

The CONFIG-MSU code has been written by **Robert Vaupel - IBM (STSM, z/OS Workload and Performance Management)**


Support : Support will be provided on a "best effort" basis. Send the spreadsheet for an analysis to alain_maneville@fr.ibm.com

The TASK BAR – Will appear when opening the worksheet

The Goal : An easy way to go from worksheet to worksheet and perform actions



The Worksheets and their usage



Use Case #1 – A real customer configuration

Spreadsheet-UC1 - General Usage explanation

Use Case #2 – A real customer configuration

Spreadsheet UC2 - How to use the EXPERT spreadsheet recommendations

Use Case #3 – A configuration – with special features....

Spreadsheet UC3 - The effect of:

Accept SPECIAL ConF ?

Use Case #4 – Special calculation – 1/4

Spreadsheet UC4 - Can we have a VM with an entitlement of 100% ?

Use Case #4 – Special calculation – 2/4 – LPAR W014

Spreadsheet UC4 - Can we have a VM with an entitlement of 100% ?

ID=IBM Corp. - LPARDesign V10-T01 Current zPCR Version-9.2a - SpecCfg=YES LPAR DEFINITION (CP)															
CFG-LP-VALID?		YES	ψMachine		Shared-Pool										
Machine-type		3906-712	#PhyProc		12	10									
MSU		2 077	#LPs (non-ICF, non-DED)		16	Ded-Pool									
Total Weight		1 000	Ratio LP/PP (base)		1.60	2									
Max LPAR		85	LSPR-AVG-V2R2-MI		17 294										
			#LPARs		8										
						HD supported on 3906									
LPARNAME	WEIGHT	#LP	%SHARE (By Pool)	RESERVED	Guaranteed #PP	Wkld LSPR	MinReq#LP	Check#LP	HD-HIGH#	HD-MED#	HD-MED%	HD-LOW#	#Active LPs	#Report LPs	
↑ ↓	↑ ↓	↑ ↓	↑ ↓		↑ ↓	↑ ↓	↑ ↓		↑ ↓	↑ ↓	↑ ↓	↑ ↓	↑ ↓	↑ ↓	
W013	300	3	30.0%		3.000	High	3	OK	3	0	N/A	0	3	3	
W014	200	3	20.0%		2.000	Average	2	OK	1	1	100.0%	1	2	2	
W015	200	2	20.0%		2.000	Average	2	OK	2	0	N/A	0	2	2	
W017	140	2	14.0%		1.400	High	2	OK	0	2	70.0%	0	2	2	
W018	10	2	1.0%		0.100	High	1	OK	0	1	10.0%	1	2	1	
W019	DED	2	100.0%		2.000	Average	2	OK	2	0	N/A	0	2	2	
W020	100	2	10.0%		1.000	Average	1	OK	0	1	100.0%	1	2	1	
W021	50	2	5.0%		0.500	High	1	OK	0	1	50.0%	1	2	1	
W014	(R6-6CP) - You have a potentiel entitlement of 2-VH But as you have 1-VL, one VH is moved to the VM pool to fuel future Unparked VL, and you will have 1-VH and 1-VM@100% [+] NOTE - You have 1VM and 1VL When ALL VLs will be UnParked, each VM and VL will have an entitlement of 50%														

Configuration Validation

Delete selected LPAR



If you have a VL you MUST have a VM for the VL to join the VM Pool – Or You can NOT have VH and VLs .. You need VM

Use Case #4 – Special calculation – 3/4 – LPAR W020

Spreadsheet UC4 - Can we have a VM with an entitlement of 100% ?

ID=IBM Corp. - LPARDesign V10-T01 Current zPCR Version-9.2a - SpecCfg=YES LPAR DEFINITION (CP)															
CFG-LP-VALID?	YES		ψMachine	Shared-Pool											
Machine-type	3906-712		#PhyProc	12	10										
MSU	2 077		#LPs (non-ICF, non-DED)	16	Ded-Pool										
Total Weight	1 000		Ratio LP/PP (base)	1.60	2										
Max LPAR	85		LSPR-AVG-V2R2-MI	17 294											
			#LPARs	8											
HD supported on 3906															
LPARNAME	WEIGHT	#LP	%SHARE (By Pool)	RESERVED	Guaranteed #PP	Wkld LSPR	MinReq#LP	Check#LP	HD-HIGH#	HD-MED#	HD-MED%	HD-LOW#	#Active LPs	#Report LPs	
↑ ↓	↑ ↓	↑ ↓	↑ ↓		↑ ↓	↑ ↓	↑ ↓		↑ ↓	↑ ↓	↑ ↓	↑ ↓	↑ ↓	↑ ↓	↑ ↓
W013	300	3	30.0%		3.000	High	3	OK	3	0	N/A	0	3	3	
W014	200	3	20.0%		2.000	Average	2	OK	1	1	100.0%	1	2	2	
W015	200	2	20.0%		2.000	Average	2	OK	2	0	N/A	0	2	2	
W017	140	2	14.0%		1.400	High	2	OK	0	2	70.0%	0	2	2	
W018	10	2	1.0%		0.100	High	1	OK	0	1	10.0%	1	2	1	
W019	DED	2	100.0%		2.000	Average	2	OK	2	0	N/A	0	2	2	
W020	100	2	10.0%		1.000	Average	1	OK	0	1	100.0%	1	2	1	
W021	50	2	5.0%		0.500	High	1	OK	0	1	50.0%	1	2	1	

Configuration Validation

Delete selected LPAR



W020

(R6-6CP) - You have a potentiel entitlement of 1-VH
But as you have 1-VL, one VH is moved to the VM pool to fuel future Unparked VL, and you will have 0-VH and 1-VM@100%
[+] NOTE - You have 1VM and 1VL
When ALL VLs will be UnParked, each VM and VL will have an entitlement of 50%

Remember that the 1st VL of a 2 LP configuration is ALWAYS **UnParked**

Use Case #4 – Special calculation – 4/4 – LPAR W020

Spreadsheet UC4 - Can we have a VM with an entitlement of 100% ?

- The effect of changing W014 #LP from 3 to 2

LPARNAME	WEIGHT	#LP	%SHARE (By Pool)	RESERVED	Guaranteed #PP	Wkld LSPR	MinReq#LP	Check#LP	HD-HIGH#	HD-MED#	HD-MED%	HD-LOW#	#Active LPs	#Report LPs
↑ ↓	↑ ↓	↑ ↓	↑ ↓		↑ ↓	↑ ↓	↑ ↓		↑ ↓	↑ ↓	↑ ↓	↑ ↓	↑ ↓	↑ ↓
W013	300	3	30.0%		3.000	High	3	OK	3	0	N/A	0	3	3
W014	200	3	20.0%		2.000	Average	2	OK	1	1	100.0%	1	2	2
W015	200	2	20.0%		2.000	Average	2	OK	2	0	N/A	0	2	2
W017	140	2	14.0%		1.400	High	2	OK	0	2	70.0%	0	2	2
W018	10	2	1.0%		0.100	High	1	OK	0	1	10.0%	1	2	1
W019	DED	2	100.0%		2.000	Average	2	OK	2	0	N/A	0	2	2
W020	100	2	10.0%		1.000	Average	1	OK	0	1	100.0%	1	2	1
W021	50	2	5.0%		0.500	High	1	OK	0	1	50.0%	1	2	1

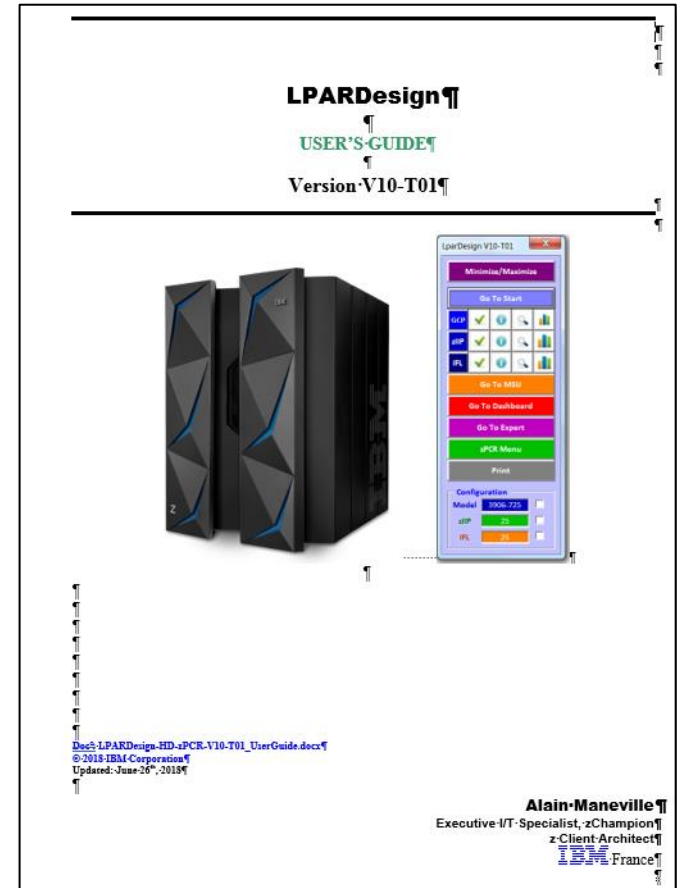
LPARNAME	WEIGHT	#LP	%SHARE (By Pool)	RESERVED	Guaranteed #PP	Wkld LSPR	MinReq#LP	Check#LP	HD-HIGH#	HD-MED#	HD-MED%	HD-LOW#	#Active LPs	#Report LPs
↑ ↓	↑ ↓	↑ ↓	↑ ↓		↑ ↓	↑ ↓	↑ ↓		↑ ↓	↑ ↓	↑ ↓	↑ ↓	↑ ↓	↑ ↓
W013	300	3	30.0%		3.000	High	3	OK	3	0	N/A	0	3	3
W014	200	2	20.0%		2.000	Average	2	OK	2	0	N/A	0	2	2
W015	200	2	20.0%		2.000	Average	2	OK	2	0	N/A	0	2	2
W017	140	2	14.0%		1.400	High	2	OK	0	2	70.0%	0	2	2
W018	10	2	1.0%		0.100	High	1	OK	0	1	10.0%	1	2	1
W019	DED	2	100.0%		2.000	Average	2	OK	2	0	N/A	0	2	2
W020	100	2	10.0%		1.000	Average	1	OK	0	1	100.0%	1	2	1
W021	50	2	5.0%		0.500	High	1	OK	0	1	50.0%	1	2	1

PLEASE ...

READ THE USER'S GUIDE

Thierry Deleris will now explain

- The zPCR Interface
- The LPARDesign Extended Tool





Merci !

Alain Maneville

